

Proposal Summary (Submission Draft)

Base repository URL (replace if needed): - <https://github.com/petroslamb/eth-llm-poc/blob/main>

This document is the submission-ready proposal summary for the Ethereum Foundation RFP on integrating LLMs into protocol security research. It is grounded in a working system (PoC 5) and describes a low-risk, auditable path to full execution and consensus coverage in 4–6 months.

Executive Summary

We propose a deterministic, auditable LLM-assisted verification system that ingests Ethereum specifications, maps EIP obligations to spec and client code, and produces reproducible discrepancy reports. The system already exists as PoC 5 for the execution layer and runs in CI with structured artifacts. We intentionally avoided complex multi-agent architectures and RAG because they reduce reproducibility, increase failure modes, and complicate audit trails. Instead, we use direct chained phases with strict boundaries, a proven production agent, and versioned artifacts.

The plan scales to consensus specs and clients, adds cross-client validation, and introduces quality thresholds and dashboards as optional extensions. The outcome is a practical, scalable verification workflow that EF can run today and extend over time. The core risk is already retired because the pipeline is live; the remaining work is coverage expansion and operational hardening.

Key Differentiators

- **Working PoC today:** execution-specs pipeline already runs with reusable CI workflows and auditable artifacts.
 - **Deterministic and auditable:** strict phase boundaries, versioned runs, and structured outputs minimize ambiguity.
 - **Lowest-risk architecture:** avoided multi-agent and RAG complexity after evaluation; chosen approach is reproducible.
 - **Clear scaling unit:** `eip-verify` job enables batch execution by EIP × client without architecture change.
-

1. RFP Objectives and How We Meet Them

RFP Objective	How We Address It	Evidence
Automated spec compliance	Extract obligations, map to spec + client code, produce gap reports	poc5/docs/POC_IMPLEMENTATION_SPEC.md
Workflow integration	Reusable GitHub Actions workflow with auditable artifacts	poc5/docs/proposal/INTEGRATION_GUIDE.md
Efficiency and accuracy	Deterministic phases, low false positives, optional LLM judge pass	poc5/docs/proposal/EVALUATION_CRITERIA_RESPONSES.md

2. Why This Approach (and Why Now)

Ethereum's spec and client surface is broad and changes frequently. The right solution is not the most complex one, but the most reproducible and auditable. After evaluating multi-agent systems, MCP/A2A protocols, RAG, symbolic repo maps, and bleeding-edge research methods, we found they introduced instability, unpredictable results, and unnecessary integration risk. The current PoC demonstrates that a simpler chained approach produces stronger accuracy and reliable artifacts. This is the lowest-risk path to a productionizable workflow that the EF can trust.

3. Scope Boundaries

In scope (Phase 1–4): - Execution-specs ingestion and EIP obligation extraction. - Execution clients (forked in repo) validated per EIP and per fork. - Consensus-specs ingestion and consensus client validation. - Deterministic artifacts, manifests, and summary reports.

Out of scope (unless explicitly requested): - Automatic code changes or patches. - Formal verification tooling beyond structured discrepancy reports. - Production deployment inside client release pipelines.

4. System Overview (Ingest → Analyze → Report)

Ingest: EIP markdown, execution-specs, and client repos.

Analyze: chained phases (extract → locate spec → analyze spec → locate client → analyze client).

Report: CSV indices, manifests, and summaries for auditability.

Output	Purpose
<code>obligations_index.csv</code>	Spec-side obligation mapping
<code>client Obligations_index.csv</code>	Client-side mapping and gaps
<code>run_manifest.json</code>	Per-phase metadata for audit
<code>summary.md / summary.json</code>	Human + machine readable reports

Design principles: simplicity, determinism, auditability, reproducibility.

5. Technical Approach (Methodologies, Frameworks, Tools)

Methodology: direct chained agent calls with strict phase boundaries and deterministic artifacts.

Rejected approaches: loosely coupled multi-agent systems, MCP/A2A protocols, RAG, symbolic repo maps, and bleeding-edge research methods due to complexity and poor reproducibility.

Frameworks/models: Claude Agent SDK with Opus 4.5 as primary, Sonnet 4.5 as fallback; evaluated GPT-5.2, Gemini 3 Pro, LangChain deep agent, and aider repomap.

Tools: native filesystem and CLI tools with structured outputs and manifest metadata.

6. Deliverables and Acceptance Criteria

Deliverable	Acceptance Criteria	Evidence
Technical architecture & design	Architecture + approach cover ingest, analysis, report, and toolchain	poc5/docs/proposal/TECHNICAL_ARCHITECTURE_AND_DESIGN.md
Working prototype	PoC 5 runs per-EIP/per-client pipeline with artifacts	poc5/docs/POC_IMPLEMENTATION_SPEC.md
Integration guidelines	Reusable workflow integration documented	poc5/docs/proposal/INTEGRATION_GUIDE.md
Operations & extension	Setup, maintenance, and future phases documented	poc5/docs/proposal/OPERATIONS_AND_EXTENSION.md

7. Success Metrics (Initial Targets)

Targets are refined with EF in Phase 1. Initial targets reflect feasible, auditable outcomes for a 4–6 month roadmap.

Metric	Initial Target	Evidence/Method
Coverage	100% of selected EIPs per fork mapped to spec obligations	CSV indices + manifests
Client validation	Per-EIP runs across agreed execution and consensus clients	Run manifests + summaries
Accuracy	<=5% false positives after judge pass on sampled runs	Cross-model review + audit logs
Reproducibility	100% artifact completeness per run	Manifests + structured outputs
Throughput	200 runs/month baseline; batch support by EIP × client	CI batch runs
Run time	Target <=60 minutes per CI run on baseline runners	CI run logs

8. Project Plan and Timeline (4–6 Months)

Phase	Timing	Outputs
Phase 1	Month 1	Validate pipeline, harden phase separation, tighten prompts
Phase 2	Month 2	Execution clients matrix coverage, repeatable batch runs
Phase 3	Month 3	Consensus-specs ingestion and obligation extraction
Phase 4	Month 4	Consensus clients matrix, EL/CL linkage
Phase 5 (optional)	Month 5	CI gating, quality thresholds, dashboarding

Phase	Timing	Outputs
Phase 6 (optional)	Month 6	Extended phases for broader protocol security mapping

Detailed plan: poc5/docs/proposal/PROJECT_PLAN_AND_TIMELINE.md

Optional extension goals (Month 5–6): - Cross-layer invariants and EL/CL consistency checks. - EIP drift detection across forks and client versions. - Higher-level protocol security mapping goals as prioritized with EF.

9. Evaluation Criteria (RFP)

Criterion	PoC Evidence	Future Expansion
Scalability	Single EIP or batch; <code>eip-verify</code> unit scales across clients	Extend to consensus specs/clients and additional phases
Accuracy	Opus 4.5 yields minimal false positives; judge pass optional	Tune prompts, add evaluators, expand validation
Reliability	Simple chained pipeline with deterministic outputs	Add harnesses, logs, monitoring
Security	Runs inside CI; report-only outputs; minimal surface	Add tighter sandboxing as needed

10. Risks and Mitigations

Risk	Mitigation
Spec ambiguity or EIP overlap	Obligation extraction with citations and audit trails
Model drift or regressions	Fixed prompts, versioned runs, optional judge pass
Cost variability	Explicit run budgeting, batch scheduling, model fallback
Client divergence	Per-client runs, EL/CL separation, artifact comparisons

Risk	Mitigation
Over-complexity	Avoid multi-agent/RAG to preserve determinism

11. Budget and Cost Structure (EUR)

We provide a cost model that separates experimentation runs from batch coverage runs, using published model pricing and a realistic CI baseline. The default estimate assumes 200 runs/month with 4M tokens/run and provides batch formulas for EIP \times client coverage. Opus is used when higher reasoning fidelity is needed; Sonnet provides a cost-effective baseline for breadth runs.

Cost Item	Estimate
Solo delivery (4 months, discounted)	EUR 53,760
Solo delivery (6 months, discounted)	EUR 80,640
LLM runs (Opus, 200/month)	EUR 6,751.20/month
LLM runs (Sonnet, 200/month)	EUR 4,050.72/month
CI (Linux baseline, 200 runs)	EUR 81.01/month

Full cost model: [poc5/docs/proposal/BUDGET_AND_COST_STRUCTURE.md](#)

CSV breakdown: [poc5/docs/proposal/BUDGET_AND_COST_STRUCTURE.csv](#)

12. Vendor Background

Petros Lambropoulos is an independent consultant with experience building production-grade agentic systems, evaluation pipelines, and ML monitoring workflows. Prior roles include AI and ML systems at Google Cloud and Workable, plus reliability and monitoring work at NannyML. This project also delivers the working PoC 5 pipeline included in the repository. References are available on request.

Docs:

[poc5/docs/proposal/VENDOR_BACKGROUND_AND_REFERENCES.md](#)

[poc5/docs/proposal/resume2025_public.pdf](#)

13. Assumptions and Dependencies

- Access to forked execution and consensus client repositories in the project.
- Ability to run GitHub Actions workflows for batch jobs.

- EF feedback cadence on scope selection and acceptance criteria.
-

14. Supporting Materials (Repository Links)

Append these to the base repository URL above:

- /poc5/docs/proposal/TECHNICAL_ARCHITECTURE_AND_DESCRIPTION.md
- /poc5/docs/proposal/PROJECT_PLAN_AND_TIMELINE.md
- /poc5/docs/proposal/BUDGET_AND_COST_STRUCTURE.csv
- /poc5/docs/proposal/INTEGRATION_GUIDE.md
- /poc5/docs/proposal/OPERATIONS_AND_EXTENSION.md
- /poc5/docs/proposal/EVALUATION_CRITERIA_REQUIREMENTS.md
- /poc5/docs/proposal/VENDOR_BACKGROUND_AND_REFERENCES.md
- /poc5/docs/proposal/PROPOSAL_READINESS_ASSESSMENT.md
- /poc5/docs/proposal/Request_for_Proposal_(RFP)_Integrating_Large_Language_Models_(LLMs)_into_Ethereum_Proto

col_Security_Research.md