

Proposal Summary

Base repository URL: <https://github.com/petroslamb/eth-llm-poc>

This document is the submission-ready proposal summary for the Ethereum Foundation RFP on integrating LLMs into protocol security research. It is grounded in a working system delivered in the eth-llm-poc repository, with the installable Python package `eip-verify`. It describes a low-risk, auditable path to full execution and consensus coverage in 4-6 months.

Executive Summary

Ethereum's Protocol Security team manually audits client implementations against evolving specifications. **eth-llm-poc** automates obligation extraction, spec mapping, and client gap detection using direct chained agent calls with strict phase boundaries. The prototype is complete: validated runs across EIP-1559, EIP-2930, and EIP-7702 with structured artifacts (CSV, JSON, Markdown) in GitHub Actions CI.

This proposal extends the prototype to production coverage. Milestones 1-2 harden the pipeline and establish accuracy baselines across an execution client matrix. Milestones 3-4 add consensus-specs ingestion and consensus client coverage. Optional Milestones 5-6 integrate CI gating and quality dashboards.

Key Differentiators

- **Working code, not a proposal:** eth-llm-poc runs end-to-end today with CI integration.
 - **Structured outputs:** CSV indices, JSON manifests, and Markdown reports provide full traceability from EIP text to client code.
 - **Validated model selection:** Claude Opus 4.5 was chosen after cross-model comparison (Haiku, Sonnet, GPT-5.2, Gemini 3 Pro); quantitative baselines are Phase 1 work.
 - **Batch-ready CLI:** `eip-verify` supports single-EIP and batch execution, scaling directly to EIP x client matrices.
-

Evidence and Validation

Validation combines qualitative evaluation (manual spot-checks and cross-model comparison) with multiple pipeline runs. Formal quantitative accuracy baselines are scoped for Phase 1.

Current validation includes manual spot-checks of obligations against EIP text, execution-specs, and Geth sources. Cross-model comparison across Haiku, Sonnet, and Opus assessed stability and quality differences. Opus 4.5 produces the most accurate outputs; Haiku showed high noise in client mapping.

Evidence sources include a qualitative evaluation summary in `docs/QUALITATIVE_EVALUATION.md`, a validation transcript in `examples/qualitative_validation_transcript.md`, example CI runs for EIP-7702 (runs 21571909617 and 21570420032), and local runs for EIP-1559 and EIP-2930 in `examples/runs/`.

1. RFP Objectives and How We Meet Them

The table below maps each RFP requirement to concrete outputs.

| RFP Objective | How We Address It | Evidence |
|---------------------------|---|------------------------------------|
| Automated spec compliance | Extract obligations, map to spec + client code, produce gap reports | Run artifacts + indexed CSVs |
| Workflow integration | Reusable GitHub Actions workflow with auditable artifacts | CI workflows + job manifests |
| Efficiency and accuracy | Phased pipeline with model selection based on observed quality | Phase manifests + evaluation notes |

2. Why This Approach

Ethereum's spec and client surface is broad and evolves frequently. The right solution is the most **reproducible and auditable**, not the most complex.

Approaches evaluated:

| Approach | Finding |
|--------------------------------------|--|
| Multi-agent systems | Coordination complexity, harder reproducibility, opaque failures |
| RAG pipelines | Context drift, retrieval failures, fragile context windows |
| Symbolic repo maps | High engineering overhead, unpredictable results |
| LangChain deep agent / aider repomap | More layers without accuracy gains |

Chosen approach: Direct chained agent calls with strict phase boundaries using Claude Agent SDK. Minimal layers between LLM and codebase preserves traceability. This is the lowest-risk path to a trustworthy workflow.

3. Scope Boundaries

In scope (Milestones 1-4): Execution-specs ingestion and EIP obligation extraction. Execution client validation starting with Geth, with a client matrix planned. Consensus-specs ingestion and consensus client validation. Deterministic artifacts, manifests, and summary reports.

Out of scope (unless explicitly requested): Automatic code changes or patches. Formal verification tooling beyond structured discrepancy reports. Production deployment inside client release pipelines.

4. System Overview (Ingest → Analyze → Report)

Pipeline flow:

1. **Ingest** → EIP markdown, execution-specs, client repos
2. **Phase 0A** → Extract obligations from EIP
3. **Phase 1A/1B** → Locate and analyze spec implementations
4. **Phase 2A/2B** → Locate and analyze client implementations
5. **Report** → CSV indices, manifests, summary reports

| Output | Purpose |
|---|----------------------------------|
| <code>obligations_index.csv</code> | Spec-side obligation mapping |
| <code>client_obligations_index.csv</code> | Client-side mapping and gaps |
| <code>run_manifest.json</code> | Per-phase metadata for audit |
| <code>summary.md / summary.json</code> | Human + machine readable reports |

Design principles: simplicity, auditability, reproducibility.

5. Technical Approach (Methodologies, Frameworks, Tools)

Methodology: Direct chained agent calls with strict phase boundaries and deterministic artifacts.

Frameworks and models:

| Category | Decision | Rationale |
|-------------------------|---------------------------------------|---|
| Primary agent | Claude Agent SDK | Best performance and reliability observed |
| Primary model | Claude Opus 4.5 | Highest accuracy in validation |
| Fallback model | Claude Sonnet 4.5 | Cost-effective for batch runs |
| Other models tested | GPT-5.2, Gemini 3 Pro | Lower quality for this task |
| Agent frameworks tested | LangChain Deep Agent, Aider + RepoMap | More complexity, weaker results |

Tools: Native filesystem and CLI tools with structured outputs and manifest metadata.

6. Deliverables and Acceptance Criteria

Each deliverable has a clear acceptance criterion so EF can validate progress without ambiguity.

| Deliverable | Acceptance Criteria | Evidence |
|---------------------------------|---|------------------------------------|
| Technical architecture & design | Architecture covers ingest, analysis, report, and toolchain | Architecture doc + system diagrams |
| Working prototype | eth-llm-poc runs per-EIP pipeline on Geth with artifacts | CLI pipeline + run outputs |
| Integration guidelines | Reusable workflow integration documented | Workflow usage + examples |
| Operations & extension | Setup, maintenance, and future phases documented | Ops guide + extension plan |

7. Success Metrics (Initial Targets)

Targets are refined with EF in Milestone 1. Current targets reflect feasible outcomes for a 4-6 month roadmap.

| Metric | Initial Target | Current State | Measurement |
|-----------------|---|---|--|
| Coverage | 100% of selected EIPs per fork mapped | EIP-1559, EIP-2930, EIP-7702 validated | CSV indices + manifests |
| Accuracy | <=5% false positives after Milestone 1 tuning | Qualitative: Opus shows low observed FP | Ground truth dataset + precision/recall in Milestone 1 |
| Reproducibility | 100% artifact completeness per run | Achieved in current runs | Manifests + structured outputs |
| Throughput | 200 runs/month baseline | CI batch runs functional | Batch workflow logs |
| Run time | <=60 minutes per CI run | ~30 min observed for EIP-7702 | CI run logs |

Accuracy note: Current validation is qualitative (spot-checks, cross-model comparison). Milestone 1 establishes quantitative baselines with curated ground truth for 2-3 well-understood EIPs.

8. Project Milestones and Timeline (4-6 Months)

The plan expands coverage in a controlled way: first harden the pipeline and establish accuracy baselines, then scale across execution and consensus layers.

| Milestone | Timing | Dependencies | Outputs |
|---------------|---------|-------------------|---|
| M0 (complete) | Done | None | Working CLI, reusable workflow, run artifacts |
| M1 | Month 1 | None | Ground truth dataset, accuracy baselines, prompt tuning |
| M2 | Month 2 | M1 accuracy >=80% | Execution client matrix (3+ clients), batch coverage |
| M3 | Month 3 | Parallel to M2 | Consensus-specs ingestion, obligation extraction |

| Milestone | Timing | Dependencies | Outputs |
|---------------|---------|--------------|---|
| M4 | Month 4 | M2 and M3 | Consensus client matrix, EL/CL linkage |
| M5 (optional) | Month 5 | None | CI gating, quality thresholds, dashboarding |
| M6 (optional) | Month 6 | None | Extended milestones for broader protocol security mapping |

Critical path: M1 accuracy validation gates M2 expansion.

Contingency: M5-6 absorb schedule slip from core milestones if needed.

Detailed plan: see Supporting Materials (Project plan and timeline).

9. Evaluation Criteria (RFP)

| Criterion | Evidence | Future Expansion |
|-------------|--|--|
| Scalability | Single EIP or batch across a fork; <code>eip-verify</code> scales in CI | Client matrices, parallel runs |
| Accuracy | Opus 4.5 yields validated outputs; quantitative baselines in Milestone 1 | Prompt tuning, evaluators, expanded validation |
| Reliability | Simple chained pipeline with deterministic outputs | Harnesses, logs, monitoring |
| Security | Runs inside CI; report-only outputs; minimal surface | Tighter sandboxing as needed |

10. Risks and Mitigations

We intentionally surface real failure modes observed during development.

| Risk | Observed Evidence | Mitigation |
|--|---|--|
| LLM hallucination (confident but incorrect mappings) | Haiku runs showed high noise in client locations (many ABI/test file false positives) | Model selection (Opus primary), cross-model review, ground truth regression |
| Spec ambiguity (multiple valid interpretations) | Some obligations (e.g., EIP1559-OBL-030) appear questionable relative to spec text | Citation-based extraction, manual arbitration workflow, disputed cases flagged |
| Extraction incompleteness (missing obligations) | Opus runs occasionally missed constraint details | Cross-EIP comparison, manual review of edge cases, iterative prompt tuning |
| Model drift (provider updates change behavior) | Not yet observed | Pinned model versions, regression suite, phased rollout of updates |
| Cost overruns | Token usage varies by EIP complexity | Budget monitoring, <code>--max-turns</code> limits, Sonnet fallback for batch runs |
| Client variation (patterns LLM doesn't recognize) | Different file structures across clients | Per-client prompt tuning, constrain to core paths, false negative tracking |

11. Budget and Cost Structure (EUR)

We provide a cost model separating engineering effort from operational costs. Opus is used for high-fidelity runs; Sonnet for cost-effective batch coverage.

| Cost Item | Estimate |
|--------------------------------------|-----------------|
| Solo delivery (4 months, discounted) | EUR 53,760 |
| Solo delivery (6 months, discounted) | EUR 80,640 |
| LLM runs (Opus, 200/month) | EUR 6,751/month |
| LLM runs (Sonnet, 200/month) | EUR 4,051/month |
| CI (Linux baseline, 200 runs) | EUR 81/month |

Cost controls: - `--llm-mode fake` for zero-cost CI and regression runs. - `--max-turns` to limit conversation length. - Phase selection to avoid unnecessary model calls.

Payment terms: - LLM and CI operational costs: invoiced upfront at project start. - Engineering delivery: invoiced monthly, or per milestone if timelines shift.

Full cost model: see Supporting Materials (Budget and cost structure).

CSV breakdown: see Supporting Materials (Budget CSV).

12. Vendor Background

Petros Lambropoulos is an independent consultant with 13 years of experience in software engineering, ML systems, and production-grade AI.

Career highlights:

- **Workable (2016-2019):** Senior Software Engineer on the NLP team. Built resume parsing and job matching systems.
- **NannyML (2021-2023):** Senior Software Engineer. Built ML monitoring platform for model drift detection.
- **Recent consulting:** Hedera/CNO (compliance-first tokenization infrastructure), dikaio.ai (agentic workflows and evaluation pipelines).

Vendor links:

- Website: <https://petroslamb.github.io/peterlamb/>
- Blog: <https://lambpetros.substack.com/>
- GitHub profile: <https://github.com/petroslamb>
- GitHub repo: <https://github.com/petroslamb/eth-llm-poc>
- LinkedIn: <https://uk.linkedin.com/in/petroslamb>

Docs: see Supporting Materials (Vendor background) and resume.

13. Assumptions and Dependencies

- Access to forked execution and consensus client repositories in the project.
 - Ability to run GitHub Actions workflows for batch jobs.
 - EF feedback cadence on scope selection and acceptance criteria.
-

14. Supporting Materials (Repository Links)

Append the folder and file below to the base repository URL above.

- Technical architecture
Folder: docs/proposal/
File: TECHNICAL_ARCHITECTURE_AND DESIGN.md

- Project plan and timeline
Folder: docs/proposal/
File: PROJECT_PLAN_AND_TIMELINE.md
- Budget and cost structure
Folder: docs/proposal/
File: BUDGET_AND_COST_STRUCTURE.md
- Budget CSV
Folder: docs/proposal/
File: BUDGET_AND_COST_STRUCTURE.csv
- Integration guide
Folder: docs/proposal/
File: INTEGRATION_GUIDE.md
- Operations and extension
Folder: docs/proposal/
File: OPERATIONS_AND_EXTENSION.md
- Evaluation criteria response
Folder: docs/proposal/
File: EVALUATION_CRITERIA_RESPONSE.md
- Vendor background
Folder: docs/proposal/
File: VENDOR_BACKGROUND_AND_REFERENCES.md
- Proposal readiness checklist
Folder: docs/proposal/
File: PROPOSAL_READINESS_CHECKLIST.md
- Canonical RFP
Folder: docs/proposal/
File: Request for Proposal (RFP)_ Integrating Large Language Models (LLMs) into Ethereum Protocol Security Research.md